

Aperture Steward - Podcast Cue Card

Study card · aim ~30-45s per beat · scan left, expand if host digs in

One-line pitch

Personal cognitive-load agent: treat attention as scarce, default-refuse noisy automation, answer as decisions with tradeoffs, keep operator-owned receipts (trace + digests). ElizaOS on Nosana; not a task butler.

Aperture metaphor: widen for breadth, stop down for depth - intentional focus control.

Arc map (host flow)

Origin -> Load problem -> Product -> ElizaOS -> Nosana -> Future of cognitive agents

1. Into agents - why cognition, not task automation?

- Agents got real; demos optimize activity (more tasks/channels).
- Wrong scarce resource. Working memory is tiny; switches leave residue; decisions degrade under churn.
- Wanted a steward, not a butler: decide/defer, protect depth of thought.

2. AI-enhanced full-stack - meaning in practice

- Model in the loop of judgment, not only autocomplete.
- AI: explore + draft · Me: stance, architecture, what we refuse.
- Own the stack end-to-end: plugin, UI, Vercel front door, Nosana agent + Qwen, correct routes (`/aperture/...`).

3. What is cognitive load - why an agent problem?

- Load = pressure on limited working memory (Cognitive Load Theory).
- Productivity tools often add load (alerts, more state). Task-only agents can multiply noise.
- Right objective: reduce selection pressure, batch, crystallize decisions, refuse low-signal automation.

4. Personal version of the problem

- [YOUR STORY - 2 sentences] Stack "worked" on paper; attention failed.
- Half-decisions across tools; easy over-automation; no surface for *what am I committing to?*
- Built a stance as software: protect attention, keep decision receipts, run on infrastructure you control.

5. How does it know load - what does it do?

- Honest: no brain sensors. Operator-declared budget/mode + character policy.
- Default-deny noisy automation; crisp plans; tradeoffs; decide/defer framing.
- Decision digests + sovereignty trace on disk you own when you log commitments.

6. What calendar / task manager can't do

Calendar / tasks | Steward

Capture + schedule | Protect attention + decision quality

More items/reminders | Refuse low-signal automation

SaaS history | Operator-owned artifacts

Calendars say *when*. Tasks say *what's open*. Steward asks *worth your working memory - and what are you committing to?*

7. Why ElizaOS - what it gave

- Challenge framework (Nosana x ElizaOS) = the arena.
- Got a real agent runtime: character, plugins, providers, actions, HTTP routes.
- Shipped stance as code: aperture plugin + steward UI, not a disposable chat wrapper.

8. Eliza constraints / friction

- Their composition: route prefixes (`/aperture/...`), env/model conventions, stock client traps.
- Thin custom stack = faster app; swarm frameworks = more orchestration (fights "less surface").
- Constraint fitted the mission: one deliberate agent.

9. Docs didn't prepare me for...

1. Routes under plugin id - judges need `/aperture/steward` + health.
2. `OPENAI\_BASE\_URL` -> Nosana `...v1`; wrong base = fake "bad key: nosana".
3. Model IDs must exist on host (Qwen, not default gpt-4o).
4. Vercel = bridge only (`AGENT\_BASE\_URL`); models live on the agent.

10. Nosana challenge - what I built

- Personal AI on ElizaOS + Nosana compute; OpenClaw-adjacent ownership theme.

- Aperture Steward: cognitive load, digests/trace, steward UI, persistent volume, Qwen inference/embeddings, optional Vercel public door.
- Narrow on purpose: clarity > engagement.

11. Design differently next time

- Light operator signals for load (still honest, not pretend EEGs).
- One unmissable demo loop: decide -> digest -> artifacts/trace.
- Align UI chat with full tool/action path.
- UX for cold starts / decentralized latency.
- Stay narrow; sovereignty + judgment over integrations.

12. Close - if cognitive agents get good?

- Work optimizes decision quality, not busyness.
- Offload noise; keep high-stakes choice; trusted owned external memory.
- Saying no becomes legible. Depth when it matters; batch the rest.
- An aperture you control - not a feed that controls you.

Science pocket (one breath)

Working memory is small -> interruptions leave residue -> micro-decisions tax judgment -> tools that maximize activity raise load. Offload commitments to inspectable, owned records; refuse automation that burns attention without leverage.

Demo (if asked · ~60s)

Open `/aperture/steward` -> decide/defer prompt -> "record a decision digest..." -> show artifacts + trace -> flash health.

Phrases to land

- "Attention as a finite resource."
- "Default-refuse noisy automation."
- "Clarity, not engagement."
- "Stance as code you can fork, deploy, and own."
- "Protect depth of thought, keep receipts, ship judgment."

Guardrails

- Don't claim biometric load detection - policy + declaration.
- Eliza: challenge first, then value, then hard-won deploy truth.
- Personal beat: insert your real week - only empty slot.
- Avoid hustle-bro / moralizing; short paragraphs; lead with the answer.

Repo: github.com/chiku524/aperture-steward · Print: 2 pages max · Study out loud once through.